

THRESHOLD ENTERPRISE SYSTEMS

Service Management Policy

Document Id	OPS-001
Title	Service Management Policy
Department	Operations
Classification	INTERNAL
Version	1.0
Status	ACTIVE
Effective Date	2026-01-15
Document Type	POLICY
Allowed Roles	EMPLOYEE, MANAGER, ENGINEER, AI_ENGINEER, SECURITY_ENGINEER, OPERATIONS_MANAGER, ADMIN

1. Policy Purpose and Enterprise Service Philosophy

This Service Management Policy establishes the authoritative governance charter, operational lifecycle principles, service ownership standards, and support frameworks governing all IT, platform, and enterprise business services provided across Threshold Enterprise Systems ('Threshold').

As a global enterprise delivering mission-critical artificial intelligence governance platforms, distributed cloud architecture, and multi-tenant SaaS applications across India, the United States, the United Kingdom, and Singapore, operational consistency and service excellence are essential. Unstructured service delivery, ambiguous ownership models, and ad-hoc support interactions introduce operational vulnerability, client dissatisfaction, and compounding technical fragility.

The purpose of this policy is to establish an end-to-end IT Service Management (ITSM) ecosystem aligned with modern ITIL v4 and Site Reliability Engineering (SRE) paradigms. It guarantees that every enterprise service operates under explicit accountability, predictable availability targets, structured change governance, rapid incident coordination, and continuous operational improvement.

Every digital asset deployed within Threshold is conceptualized as a customer-aligned or internal capability service. Operational teams must champion transparency, deterministic procedural execution, automated observability, and customer-centric support.

2. Scope and Service Architecture Governance

This policy applies universally across all business units, operational departments, development centers, and third-party managed service providers operating within Threshold Enterprise Systems.

The scope encompasses three foundational categories of enterprise services:

1. Core Commercial Services: Customer-facing SaaS products, public REST and streaming APIs (ENG-001), AI governance evaluation engines, and real-time inference endpoints.
2. Internal Enterprise Platforms: Corporate identity directories, single sign-on (SSO) gateways, employee HRIS portals, database clusters (ENG-005), and enterprise communication fabrics.
3. Operational Infrastructure Services: Cloud compute clusters, Kubernetes meshes, continuous delivery pipelines (ENG-004), network transit fabrics, and monitoring aggregators (OPS-008).

All external service integrations, vendor software-as-a-service platforms, and contracted infrastructure providers must undergo operational compliance verification prior to production onboarding.

3. The Enterprise Service Lifecycle

Every digital and operational service within Threshold must transition through five discrete, auditable lifecycle stages:

Stage 1 - Service Strategy and Inception: Business and technical stakeholders formulate business cases, analyze market demand, define compliance envelopes, and establish target availability tiers.

Stage 2 - Service Design and Architecture: Engineering and Operations collaboratively define system architecture (ENG-007), data security safeguards (SEC-001), Service Level Objectives (OPS-005), and operational data retention rules (OPS-007). Non-functional requirements including disaster recovery failover and telemetry instrumentation must be approved during this stage.

Stage 3 - Service Transition and Release: Deployment readiness is verified through automated staging gates, code reviews (ENG-002), change management review (OPS-002), and formal deployment guides (ENG-004). Sandboxed dry runs and rollback procedures must be validated before production cutover.

Stage 4 - Service Operations: Continuous production execution, 24/7 telemetry monitoring (OPS-008), proactive incident detection, standard operating procedures execution (OPS-003), and operational incident escalation (OPS-006).

Stage 5 - Continual Service Improvement (CSI): Post-incident learning, service level reviews, capacity adjustments, and systematic architecture hardening.

4. Explicit Service Ownership, RACI Matrix, and Operational Accountability

To eliminate organizational ambiguity, every deployed enterprise service must maintain an authoritative Service Record in the enterprise Configuration Management Database (CMDB) containing designated human roles:

- Service Owner (Business & Product): Accountable for business outcomes, feature roadmap, compliance alignment, customer communication, and commercial viability of the service.
- Technical Lead (Engineering): Accountable for software architecture, code quality, technical debt remediation, unit/integration test coverage, and API integration standards.
- Operational Lead (Operations / SRE): Accountable for 24/7 service health, monitoring coverage, incident runbook maintenance, capacity planning, and deployment execution.

The Enterprise Service RACI Matrix mandates:

- Responsible: Assigned SREs and engineering squad members executing daily operational maintenance and deployment pipelines.
- Accountable: The Service Owner holds singular, non-delegable accountability for service delivery integrity.
- Consulted: Information Security (SEC-001), Enterprise Governance (AIG-001), and Legal counsel on compliance and architectural modifications.
- Informed: Affected internal department users and customer support teams during scheduled maintenance or incident degradations.

Services lacking assigned, active owners in the CMDB are deemed orphaned and subject to immediate administrative review and potential decommissioning within thirty (30) days.

5. Service Availability Targets and Operational Responsibilities

Threshold categorizes all enterprise services into three distinct Operational Availability Tiers:

- Tier 1 - Mission-Critical Core Services: Revenue-generating customer endpoints, public APIs, and enterprise identity federation. Availability Target: 99.9% uptime (maximum 43.8 minutes of unplanned downtime per month). Requires 24/7 SRE active monitoring and 15-minute response SLAs.
- Tier 2 - Business-Operational Services: Internal business workflow tools, analytics pipelines, and staging environments. Availability Target: 99.5% uptime (maximum 3.65 hours of unplanned downtime per month). Requires 16/5 operational

monitoring and 1-hour response SLAs.

- Tier 3 - Non-Critical Auxiliary Services: Internal collaboration sandboxes and historical archival tools. Availability Target: 99.0% uptime during core regional business hours.

Operational responsibilities include establishing automated health checks, maintaining validated SOPs (OPS-003), conducting quarterly disaster recovery drills, and participating in on-call escalations.

Scheduled maintenance windows authorized under OPS-002 do not count against unplanned downtime calculations, provided they adhere strictly to pre-announced maintenance schedules.

6. Service Request Management and IT Support Workflows

Workforce members and enterprise stakeholders interact with Operations through the centralized Enterprise Service Desk portal:

- Standard Service Requests: Routine administrative operations (e.g., standard hardware provisioning, VPN configuration, software license assignment) execute via automated self-service workflows with standard 24-hour turnaround.
- Privileged Access Requests: Requests for elevated administrative roles or database inspection follow multi-tier approvals governed under SEC-002 and ENG-005.
- Emergency Support Ingestion: Critical service degradation tickets automatically route to the Operations Center triage desk, triggering immediate escalation bridges under OPS-006.
- Service Request SLAs: Standard requests carry strict fulfillment targets: Priority 1 (Urgent Business Blocker): ≤ 2 hours; Priority 2 (Standard Work Task): ≤ 8 business hours; Priority 3 (General Inquiry): ≤ 24 business hours.
- Automated Ticket Triage: Machine learning classifiers categorize incoming tickets, route tasks to appropriate regional support queues (Bangalore, London, San Francisco, Singapore), and flag recurring hardware or software anomalies.

7. Incident and Change Coordination Interfaces

Service management relies on strict operational separation and seamless coordination between incident handling and system change:

- Incident Coordination Interface: Operational disruptions are triaged by Operations under OPS-006, coordinating cross-functional mobilization across Engineering (ENG-006), Cybersecurity (SEC-004), and AI Governance (AIG-009). Operations ensures that incident communication bridges remain orderly and updates the public status page.
- Change Coordination Interface: All modifications to production environments must be vetted under OPS-002, ensuring that scheduled releases do not collide, violate blackout freezes, or compromise service dependencies.
- Change-Incident Feedback Loop: If an incident is linked to a recent deployment, the change record is flagged, triggering an immediate post-implementation review to prevent recurring regressions.

8. Service Performance Governance, CSI Register, and Reviews

Operations leadership conducts monthly Service Performance Reviews evaluating:

1. Service Level Performance: Adherence to agreed SLOs and error budget consumption under OPS-005.
2. Telemetry and Alarm Accuracy: Reviewing false-positive rates and monitoring coverage under OPS-008.
3. Incident Trends and Root Causes: Aggregating blameless post-mortem insights from ENG-006 and OPS-006 to eliminate recurring operational bottlenecks.
4. Business Continuity Readiness: Validating recovery time objectives (RTO) and recovery point objectives (RPO) against the Business Continuity Policy (OPS-004).

Continual Service Improvement (CSI) Register: Every identified operational enhancement, monitoring improvement, or technical debt remediation item must be cataloged in the centralized CSI register. Items are prioritized by business impact

and reviewed quarterly by the Operations Steering Committee.

9. Compliance, Exceptions, and Cross-Document Governance

Adherence to service management standards is mandatory across all global facilities. Temporary exceptions—such as deferring a required monitoring upgrade—require written authorization from the Director of Global Operations and the VP of Engineering.

Cross-Document Operational Governance Interconnections:

- OPS-002: Change Management Policy (Governs release authorization and change risk assessment).
- OPS-003: Standard Operating Procedures Policy (Governs procedural runbook creation).
- OPS-004: Business Continuity Policy (Establishes disaster recovery and continuity plans).
- OPS-005: Service Level Management Policy (Defines SLOs, SLAs, and error budgets).
- OPS-006: Operational Incident Escalation Procedure (Defines multi-domain incident routing).
- OPS-007: Operational Data Retention Guidelines (Regulates log and telemetry preservation).
- OPS-008: Service Performance Monitoring Guidelines (Governs telemetry and alerting baselines).